

V Namratha  
1NT19IS180  
C2 batch

```
package nam;
```

```
import java.io.IOException;  
import java.util.Iterator;  
import java.util.StringTokenizer;
```

```
import org.apache.hadoop.fs.Path;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.LongWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapred.FileInputFormat;  
import org.apache.hadoop.mapred.FileOutputFormat;  
import org.apache.hadoop.mapred.JobClient;  
import org.apache.hadoop.mapred.JobConf;  
import org.apache.hadoop.mapred.MapReduceBase;  
import org.apache.hadoop.mapred.Mapper;  
import org.apache.hadoop.mapred.OutputCollector;  
import org.apache.hadoop.mapred.Reducer;  
import org.apache.hadoop.mapred.Reporter;  
import org.apache.hadoop.mapred.TextInputFormat;  
import org.apache.hadoop.mapred.TextOutputFormat;
```

```
public class WordCount {
```

```
    public static class Map extends MapReduceBase implements  
        Mapper<LongWritable, Text, Text, IntWritable> {  
        private final static IntWritable one = new IntWritable(1);  
        private Text word = new Text();  
        public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable>  
            output, Reporter reporter) throws IOException {  
            String line = value.toString();  
            StringTokenizer tokenizer = new StringTokenizer(line);  
            while (tokenizer.hasMoreTokens()) {  
                word.set(tokenizer.nextToken());  
                output.collect(word, one);  
            }  
        }  
    }
```

```
}  
}
```

```
public static class Reduce extends MapReduceBase implements  
Reducer<Text, IntWritable, Text, IntWritable> {  
    public void reduce(Text key, Iterator<IntWritable> values,  
        OutputCollector<Text, IntWritable> output, Reporter reporter)  
        throws IOException {  
        int sum = 0;  
        while (values.hasNext()) {  
            sum += values.next().get();  
        }  
        output.collect(key, new IntWritable(sum));  
    }  
}
```

```
public static void main(String[] args) throws Exception  
{  
    JobConf conf = new JobConf(WordCount.class);  
    conf.setJobName("wordcount");  
    conf.setOutputKeyClass(Text.class);  
    conf.setOutputValueClass(IntWritable.class);  
    conf.setMapperClass(Map.class);  
    conf.setCombinerClass(Reduce.class);  
    conf.setReducerClass(Reduce.class);  
    conf.setInputFormat(TextInputFormat.class);  
    conf.setOutputFormat(TextOutputFormat.class);  
    FileInputFormat.setInputPaths(conf, new Path(args[0]));  
    FileOutputFormat.setOutputPath(conf, new Path(args[1]));  
    JobClient.runJob(conf);  

```

```
}
```

```
}
```

```
Activities Terminal Jun 18 12:20
hadoop@admin1-HP-280-G4-MT-Business-PC: ~/hadoop-3.2.1/sbin

hadoop@admin1-HP-280-G4-MT-Business-PC:~$ cd $HADOOP_HOME/sbin
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [admin1-HP-280-G4-MT-Business-PC]
Starting resourcemanager
Starting nodemanagers
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ jps
5474 SecondaryNameNode
5235 DataNode
2101 org.eclipse.equinox.launcher_1.5.600.v20191014-2022.jar
6165 Jps
5078 NameNode
5814 NodeManager
5662 ResourceManager
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -mkdir -p /input180
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -appendToFile - /input180 /text.txt
appendToFile: stdin (-) must be the sole input argument when present
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -appendToFile - /input180/text.txt
ghe who must not be named
hhe is voldermort2022-06-18 12:16:51,080 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hadoop jar /home/hadoop/Documents/namratha.jar /input180 /output180
2022-06-18 12:18:38,126 INFO client.RMProxy: Connecting to ResourceManager at /127.0.0.1:8032
2022-06-18 12:18:38,584 INFO client.RMProxy: Connecting to ResourceManager at /127.0.0.1:8032
2022-06-18 12:18:38,750 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2022-06-18 12:18:39,030 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1655534645081_0001
2022-06-18 12:18:39,215 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2022-06-18 12:18:39,333 INFO mapred.FileInputFormat: Total input files to process : 1
2022-06-18 12:18:39,368 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2022-06-18 12:18:39,417 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
```

```
Activities Terminal Jun 18 12:20
hadoop@admin1-HP-280-G4-MT-Business-PC: ~/hadoop-3.2.1/sbin

appendToFile: stdin (-) must be the sole input argument when present
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -appendToFile - /input180/text.txt
ghe who must not be named
hhe is voldermort2022-06-18 12:16:51,080 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hadoop jar /home/hadoop/Documents/namratha.jar /input180 /output180
2022-06-18 12:18:38,126 INFO client.RMProxy: Connecting to ResourceManager at /127.0.0.1:8032
2022-06-18 12:18:38,584 INFO client.RMProxy: Connecting to ResourceManager at /127.0.0.1:8032
2022-06-18 12:18:38,750 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2022-06-18 12:18:39,030 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1655534645081_0001
2022-06-18 12:18:39,215 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2022-06-18 12:18:39,333 INFO mapred.FileInputFormat: Total input files to process : 1
2022-06-18 12:18:39,368 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2022-06-18 12:18:39,417 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2022-06-18 12:18:39,434 INFO mapreduce.JobSubmitter: number of splits:2
2022-06-18 12:18:39,526 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2022-06-18 12:18:39,584 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1655534645081_0001
2022-06-18 12:18:39,584 INFO mapreduce.JobSubmitter: Executing with tokens: []
2022-06-18 12:18:39,882 INFO conf.Configuration: resource-types.xml not found
2022-06-18 12:18:39,882 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2022-06-18 12:18:40,084 INFO impl.YarnClientImpl: Submitted application application_1655534645081_0001
2022-06-18 12:18:40,113 INFO mapreduce.Job: The url to track the job: http://admin1-HP-280-G4-MT-Business-PC:8088/proxy/application_1655534645081_0001/
2022-06-18 12:18:40,114 INFO mapreduce.Job: Running job: job_1655534645081_0001
2022-06-18 12:18:45,229 INFO mapreduce.Job: Job job_1655534645081_0001 running in uber mode : false
2022-06-18 12:18:45,229 INFO mapreduce.Job: map 0% reduce 0%
2022-06-18 12:18:49,300 INFO mapreduce.Job: map 100% reduce 0%
2022-06-18 12:18:56,439 INFO mapreduce.Job: map 100% reduce 100%
2022-06-18 12:18:58,470 INFO mapreduce.Job: Job job_1655534645081_0001 completed successfully
2022-06-18 12:18:58,608 INFO mapreduce.Job: Counters: 54
File System Counters
FILE: Number of bytes read=104
FILE: Number of bytes written=677556
FILE: Number of read operations=0
```

```
Activities Terminal Jun 18 12:20 hdoop@admin1-HP-280-G4-MT-Business-PC: ~/hadoop-3.2.1/sbin

File System Counters
  FILE: Number of bytes read=104
  FILE: Number of bytes written=677556
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=247
  HDFS: Number of bytes written=62
  HDFS: Number of read operations=11
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
  HDFS: Number of bytes read erasure-coded=0
Job Counters
  Launched map tasks=2
  Launched reduce tasks=1
  Data-local map tasks=2
  Total time spent by all maps in occupied slots (ms)=3685
  Total time spent by all reduces in occupied slots (ms)=4117
  Total time spent by all map tasks (ms)=3685
  Total time spent by all reduce tasks (ms)=4117
  Total vcore-millisecons taken by all map tasks=3685
  Total vcore-millisecons taken by all reduce tasks=4117
  Total megabyte-millisecons taken by all map tasks=3773440
  Total megabyte-millisecons taken by all reduce tasks=4215808
Map-Reduce Framework
  Map input records=2
  Map output records=9
  Map output bytes=80
  Map output materialized bytes=110
  Input split bytes=182
  Combine input records=9
  Combine output records=9
  Reduce input groups=9
  Reduce shuffle bytes=110
  Reduce input records=9
  Reduce output records=9
```

```
Activities Terminal Jun 18 12:20 hdoop@admin1-HP-280-G4-MT-Business-PC: ~/hadoop-3.2.1/sbin

  Total vcore-millisecons taken by all map tasks=3685
  Total vcore-millisecons taken by all reduce tasks=4117
  Total megabyte-millisecons taken by all map tasks=3773440
  Total megabyte-millisecons taken by all reduce tasks=4215808
Map-Reduce Framework
  Map input records=2
  Map output records=9
  Map output bytes=80
  Map output materialized bytes=110
  Input split bytes=182
  Combine input records=9
  Combine output records=9
  Reduce input groups=9
  Reduce shuffle bytes=110
  Reduce input records=9
  Reduce output records=9
  Spilled Records=18
  Shuffled Maps =2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=160
  CPU time spent (ms)=1510
  Physical memory (bytes) snapshot=774516736
  Virtual memory (bytes) snapshot=7611191296
  Total committed heap usage (bytes)=828375040
  Peak Map Physical memory (bytes)=292495360
  Peak Map Virtual memory (bytes)=2534359040
  Peak Reduce Physical memory (bytes)=191168512
  Peak Reduce Virtual memory (bytes)=2544308224
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
```

```
Activities Terminal Jun 18 12:20
hadoop@admin1-HP-280-G4-MT-Business-PC: ~/hadoop-3.2.1/sbin

Shuffled Maps =2
Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=160
CPU time spent (ms)=1510
Physical memory (bytes) snapshot=774516736
Virtual memory (bytes) snapshot=7611191296
Total committed heap usage (bytes)=828375040
Peak Map Physical memory (bytes)=292495360
Peak Map Virtual memory (bytes)=2534359040
Peak Reduce Physical memory (bytes)=191168512
Peak Reduce Virtual memory (bytes)=2544308224
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=65
File Output Format Counters
Bytes Written=62
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -cat /output180/part*
2022-06-18 12:19:20,552 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
be 1
ghe 1
hhe 1
is 1
must 1
named 1
not 1
voldermort 1
who 1
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$
```